


Deep Implicit Layers

2021-1-7

손형욱

hyounguk.shon@kaist.ac.kr

Based on the NIPS 2020 tutorial



Deep Implicit Layers:
Neural ODEs, Equilibrium
Models, and Beyond

<http://implicit-layers-tutorial.org>

David Duvenaud
University of Toronto
and Vector Institute



J. Zico Kolter
Carnegie Mellon and
Bosch Center for AI




Matt Johnson
Google Brain



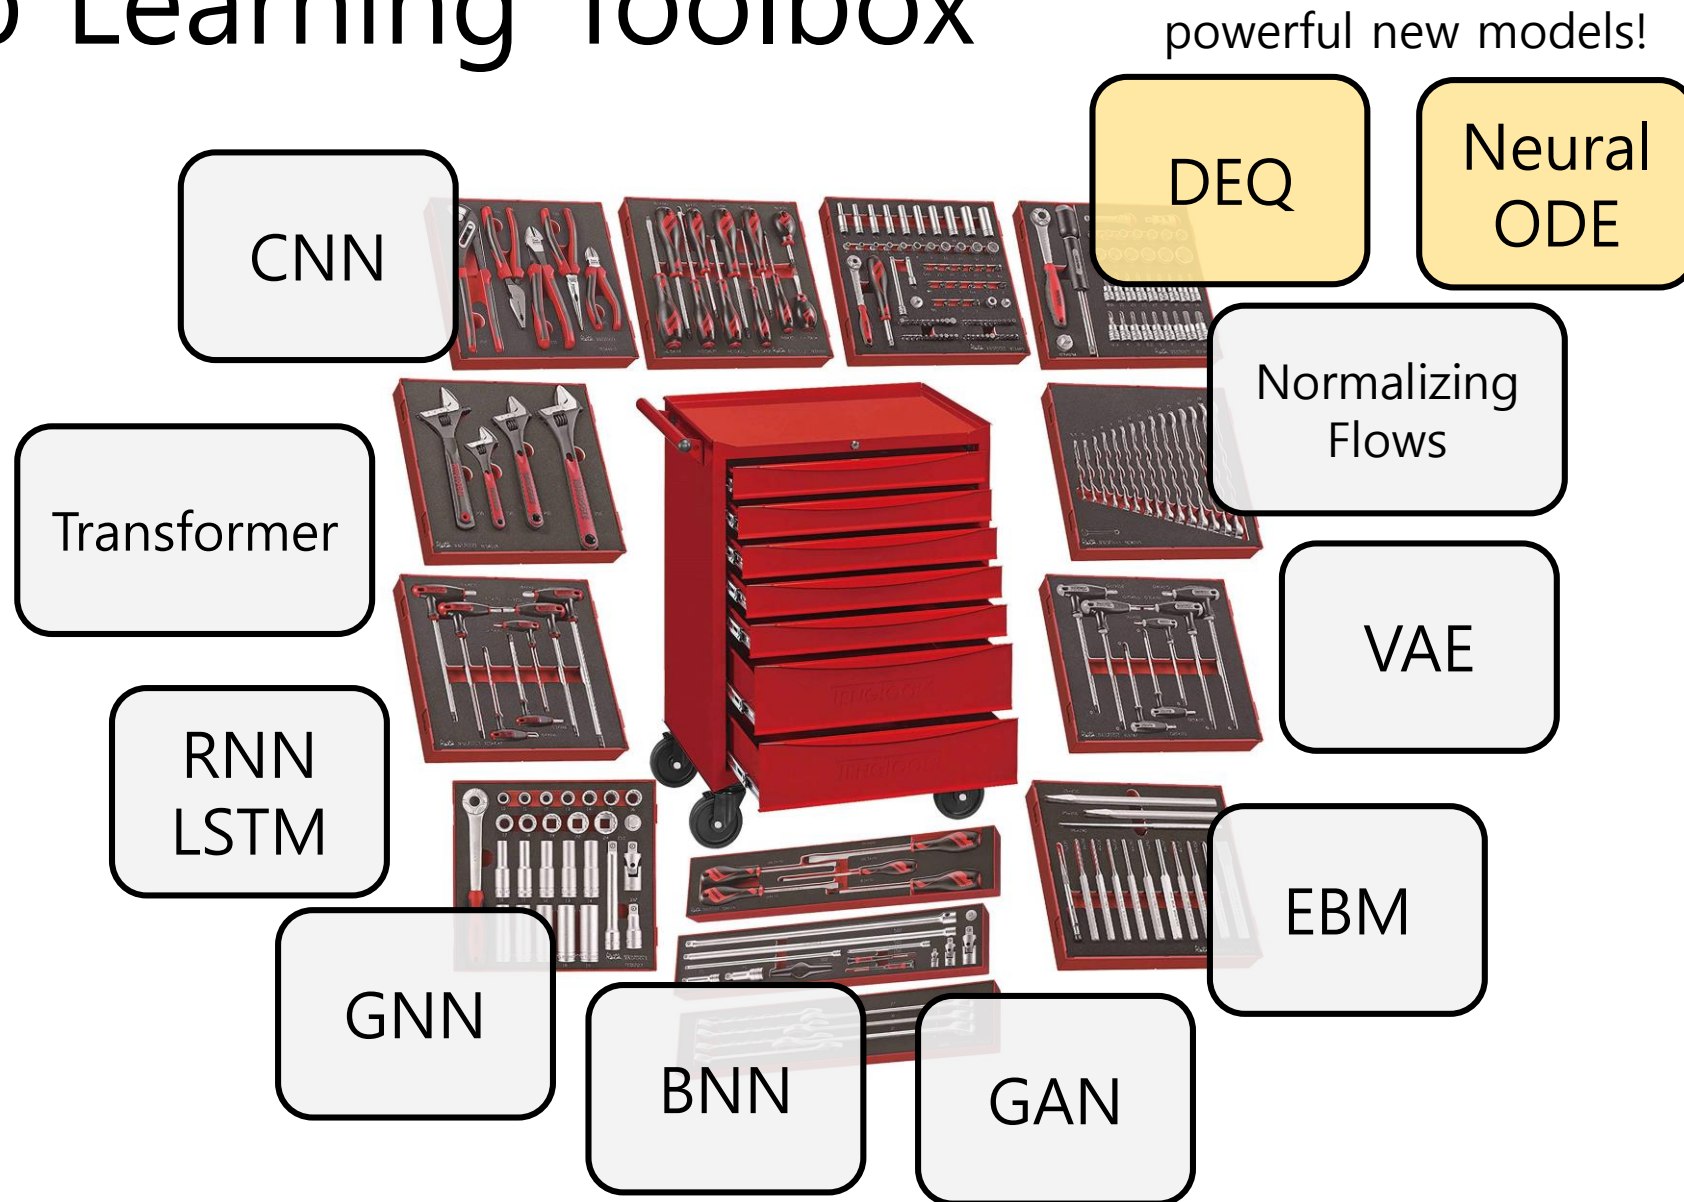
Learning implicit functions

- Implicit functions in Deep Learning:
 - Deep implicit layers
 - Implicit neural representations
- Implicit function opens a new search space for deep learning.
 - Novel architectures for learning implicit functions
 - Novel representations that simplify difficult problems
- Deep implicit layers
 - A layer is a differentiable parametric function.
 - An emerging family of neural networks. (roots back to 1980s)
- Why implicit layers?
 - More representation power
 - Extreme memory efficiency
 - Simple architecture
 - Isolates the layer's behavior from its computation.

- Occupancy Networks
- Neural Radiance Field (NeRF)



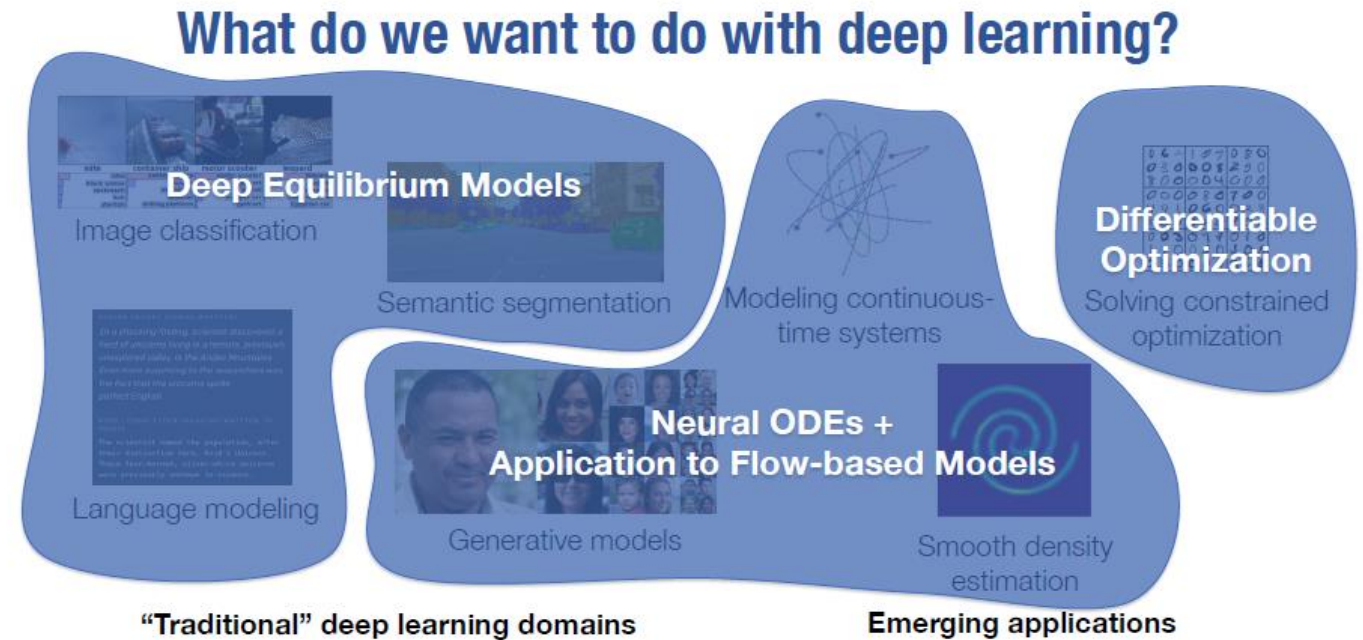
Deep Learning Toolbox



Overview of the topics

- Deep Equilibrium Models (DEQs)
- Neural Ordinary Differential Equations (Neural ODEs)
- Rules for forward and backward pass
- Applications
- Future directions

Chen et al., Neural Ordinary Differential Equations. NIPS 2018.
Bai et al., Deep Equilibrium Models. NIPS 2019.



Picture credits: [Krizhevsky et al., 2012; Bai et al., 2020; Grathwohl et al., 2018; Radford et al., 2019; Keras et al., 2018; Wang et al., 2019]

Deep Equilibrium Models


Shaojie Bai, Zico Kolter, Vladlen Koltun. *Deep Equilibrium Models*. NeurIPS 2019.

Shaojie Bai, Vladlen Koltun, Zico Kolter. *Multiscale Deep Equilibrium Models*. NeurIPS 2020.

- Examples where weight-tying was used for improving performance:
- Dehghani et al., Universal Transformers. ICLR 2019.
 - Lan et al., ALBERT: A Lite BERT for Self-supervised Learning of Language Representations. ICLR 2020.


A simple example: Fixed-Point Layer

- A deep neural network:



$$z_{i+1} = \sigma(W_i z_i + b_i)$$


- Our fixed-point network:
 1. Parameter W is shared across layers.
 2. The input x is injected to every layers.



$$z_{i+1} = \sigma(W z_i + x)$$

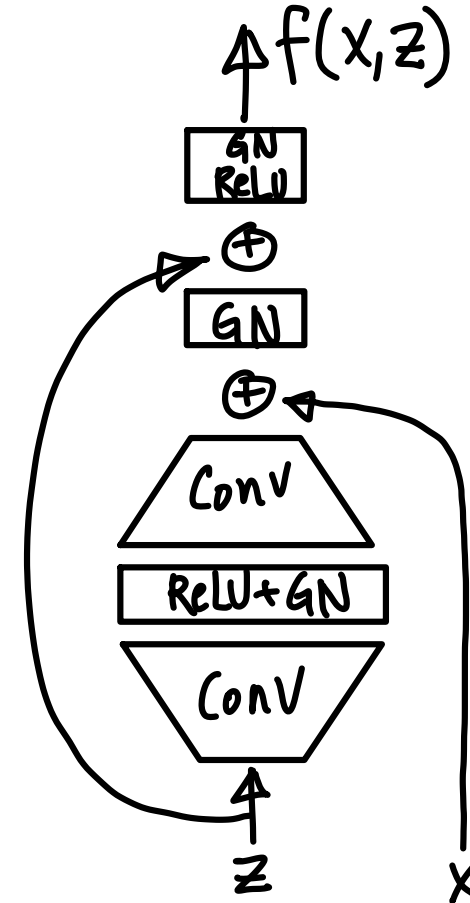
Fixed-Point Layer: forward and backward

- Fixed-Point Layer
 - We repeatedly apply the same transformation to hidden units.
 - At $k \rightarrow \infty$, the hidden unit converges to some fixed-point \mathbf{z}^* such that $\mathbf{z}^* = \sigma(W\mathbf{z}^* + \mathbf{x})$
 - We use \mathbf{z}^* as the output of our fixed-point layer.
- Forward pass is fixed-point computation
 1. Input: \mathbf{x}
 2. Iterate $\mathbf{z}_{i+1} = f(\mathbf{z}_i, \mathbf{x})$ until $\|\mathbf{z}_{i+1} - \mathbf{z}_i\| < \varepsilon$.
 3. Output: \mathbf{z}^*
- Backward pass
 - Backpropagation through time (it works, but inefficient)
- DEQ uses efficient forward and backward algorithms.
 - Fast forward computation
 - Fast and very memory-efficient backward computation



Deep Equilibrium Models

- Deep Equilibrium Models
 - So far, $f_{\theta}(\mathbf{x}, \mathbf{z}) = \sigma(W\mathbf{z} + \mathbf{x})$ was just a single layer.
 - DEQ incorporates more expressive architectures for $f_{\theta}(\mathbf{x}, \mathbf{z})$.
- Forward pass through a DEQ
 1. Input: \mathbf{x}
 2. DEQ uses fixed-point solvers that are substantially more efficient.
 3. Output: \mathbf{z}^*
- Accelerating convergence of fixed-point iterations
 - Root finding: Solve $f(\mathbf{z}, \mathbf{x}) - \mathbf{z} = \mathbf{0}$ (e.g. Broyden's method)
 - Anderson acceleration: Extrapolate \mathbf{z}_{i+1} from previous iterations.
- Anderson acceleration method:
 - $\mathbf{z}_{i+1} = \alpha_1 f(\mathbf{z}_i) + \alpha_2 f(\mathbf{z}_{i-1}) + \dots + \alpha_M f(\mathbf{z}_{i-M+1})$
 - $\sum \alpha_m = 1$
 - Choose $\boldsymbol{\alpha} = \text{argmin}[\mathbf{G}\boldsymbol{\alpha}]$, $\mathbf{G} = [f(\mathbf{z}_i) - \mathbf{z}_i \quad \dots \quad f(\mathbf{z}_{i-M+1}) - \mathbf{z}_{i-M+1}]$



A residual cell

Backward pass through a DEQ

Takeaways:


1. Only need to record the final fixed-point z^* .
2. Backward pass has only $\mathcal{O}(1)$ memory footprint.
3. Backward pass solves another fixed-point problem.

- The implicit differentiation rule

- $f(x, z) - z = 0$
- $\frac{\partial f(x, z)}{\partial x} + \frac{\partial f(x, z)}{\partial z} \frac{\partial z}{\partial x} - \frac{\partial z}{\partial x} = 0$
- $\frac{\partial z}{\partial x} = \left[I - \frac{\partial f(x, z)}{\partial z} \right]^{-1} \frac{\partial f(x, z)}{\partial x}$
- Only requires the output z^* , not the trajectory z_1, z_2, \dots, z_t .

- Backpropagation through DEQ


- We want to compute the vector-Jacobian product: $(\nabla_x \mathcal{L})^\top = (\nabla_z \mathcal{L})^\top \frac{\partial z}{\partial x}$
- $(\nabla_z \mathcal{L})^\top \frac{\partial z}{\partial x} = (\nabla_z \mathcal{L})^\top \left[I - \frac{\partial f(x, z)}{\partial z} \right]^{-1} \frac{\partial f(x, z)}{\partial x}$
- Let $v^\top = (\nabla_z \mathcal{L})^\top \left[I - \frac{\partial f(x, z)}{\partial z} \right]^{-1}$
- $v^\top = v^\top \frac{\partial f(x, z)}{\partial z} + (\nabla_z \mathcal{L})^\top$ (this is a fixed-point equation)
- Backward pass is another fixed-point problem!




*same works for the parameter gradient,
 $(\nabla_z \mathcal{L})^\top \frac{\partial z}{\partial \theta} = (\nabla_z \mathcal{L})^\top \left[I - \frac{\partial f(x, z)}{\partial z} \right]^{-1} \frac{\partial f(x, z)}{\partial \theta}$

Deep Equilibrium Models: Summary

- Inference and training DEQs
 - DEQs = infinite-depth neural networks
 - Forward pass runs a fixed-point solver.
 - Backward pass runs another fixed-point solver.
 - Fixed-points are computed efficiently via acceleration methods.
- Training DEQs (vs. traditional neural networks)
 - $\mathcal{O}(1)$ memory training regardless of depth of layers. (vs. $\mathcal{O}(L)$ memory)
 - No need to store the activations of each intermediate layer.
- Why does DEQs matter?
 - Generally performs better when model sizes are similar.
 - Requires much less GPU memory for training
 - Provides a mechanism to trade off accuracy vs. latency at test-time
- Current shortcomings
 - Training DEQs typically takes 2x – 3x longer.
 - Perhaps we can regularize DEQs to be faster to solve.





Finding the fixed-point becomes harder as the model fits the dataset. (the model learns to become 'deep')



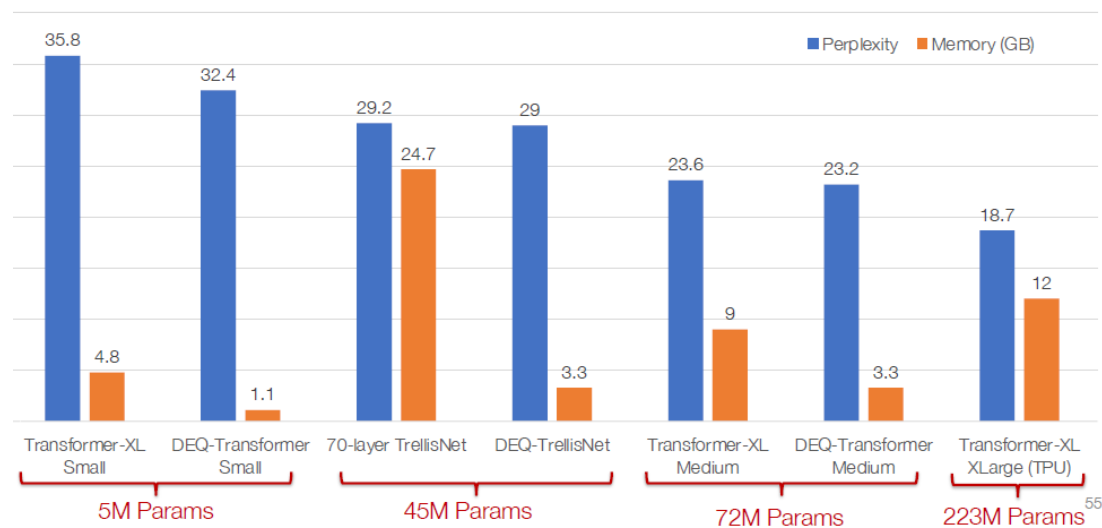
Perplexity vs. latency trade-off

Deep Equilibrium Models: applications

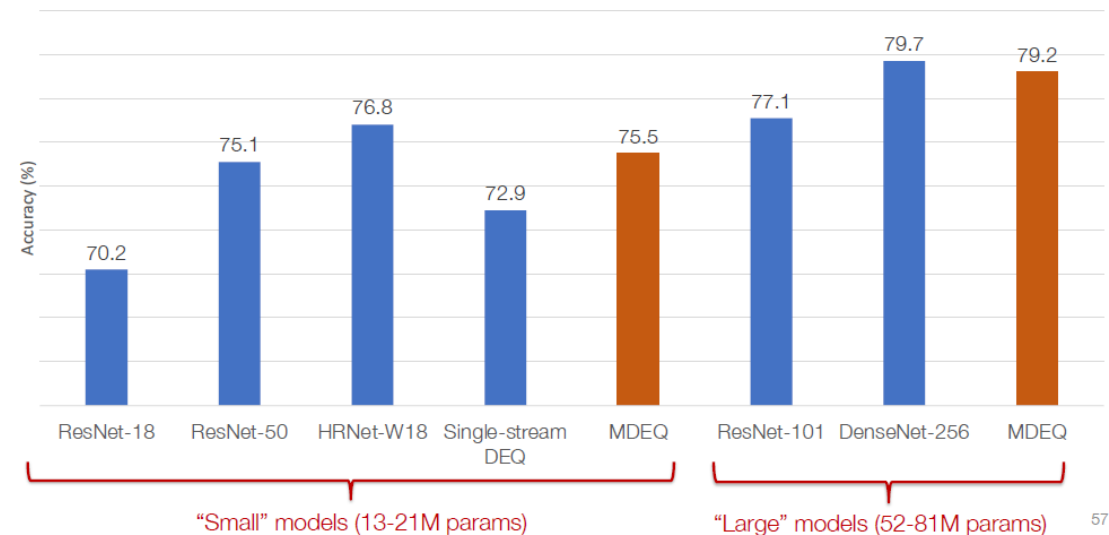
- Advantages
 - Represent modern deep networks using a single implicit layer.
 - Adaptive computation
 - Competitive performance in large-scale NLP tasks.
- Multiscale DEQ
 - Competitive performance in large-scale visual tasks.



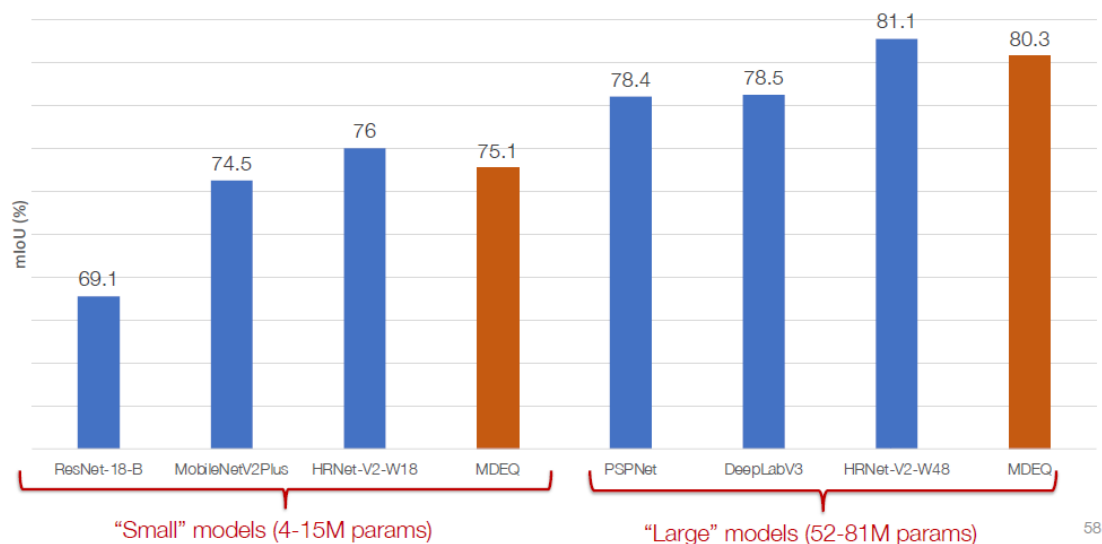
Language modeling: WikiText-103



ImageNet Top-1 Accuracy



Citiscapes mIoU






Neural ODEs

Ricky Chen, Yulia Rubanova, Jesse Bettencourt, David Duvenaud. *Neural Ordinary Differential Equations*. NeurIPS 2018.



Background: Ordinary Differential Equations

- Velocity of a free-falling object
 - ODE specifies a continuous-time state variable by its rate of change.
 - State variable: $v(t)$
 - State dynamics: $m \frac{dv}{dt} = mg - kv$
- Ordinary Differential Equations (1st order)
 - The state vector $z(t)$ follows the dynamics $f(z, t)$.
 - $$\begin{cases} \frac{dz}{dt} = f(z, t) \\ z(t_0) = z_0 \end{cases}$$
 - Dynamics + initial state = Initial value problem
- Solving an ODE means finding $z(t)$.
 - We can evaluate $z(t_1)$ by integrating $f(z, t)$ for $t \in [t_0, t_1]$.
 - $z(t_1) = z(t_0) + \int_{t_0}^{t_1} f(z, t) dt$
 - We have numerical ODE solvers. (e.g., Euler, Runge-Kutta, Dopri5)
 - Modern solvers adaptively adjust the number of function evaluations according to the complexity of the dynamics.




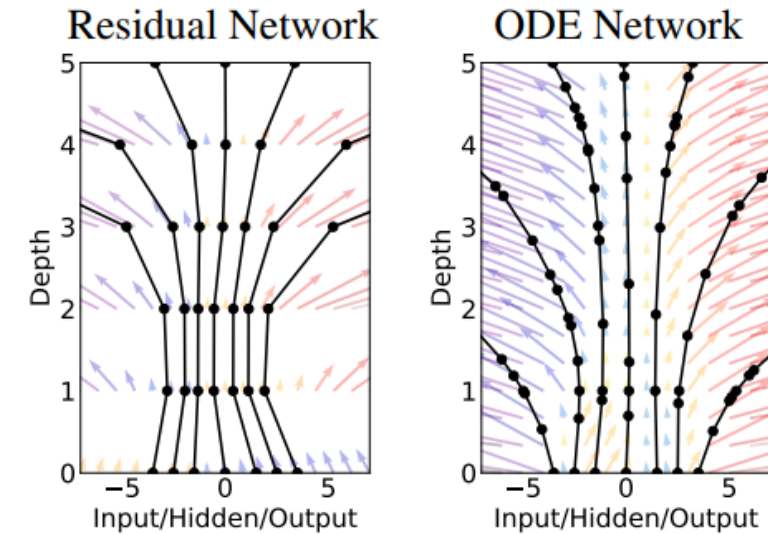
Weight-tied Residual Network

- Neural ODE is a depth-continuous ResNet.
 - Residual block: $z_{i+1} = z_i + f_{\theta}(z_i, i)$
 - Neural ODE: $dz = f_{\theta}(z, t)dt$



Neural ODE

- Parametrizes the dynamics $f(z, t)$ by a neural network.
 - $\frac{dz}{dt} = f_\theta(z, t)$
 - $z(t)$ is the hidden state vector.
 - f_θ is just a MLP. $(\mathbb{R}^{|z|+1} \rightarrow \mathbb{R}^{|z|})$
- Forward pass is an initial value problem.
 - Use numerical solver: $z(t_1) = \text{ODEint}(f_\theta, z_0, [t_0, t_1])$
 - z_0 : layer input
 - $z(t_1)$: layer output
- Neural ODE is a depth-continuous ResNet.
 - Residual block: $z_{i+1} = z_i + f_\theta(z_i, i)$
 - Neural ODE: $dz = f_\theta(z, t)dt$
- What are Neural ODEs good for?
 - Can be used anywhere a ResNet can.
 - Flexible density estimation and time-series models.
 - Whenever knowing the trajectory is important.



Dissecting Neural ODEs. Massaroli, Poli, Park, Yamashita, Asama (2020)

Backward pass through Neural ODE

- Don't backpropagate through time – there's a better way!
 - Things we want to compute: $\frac{\partial \mathcal{L}}{\partial z(0)}$ and $\frac{\partial \mathcal{L}}{\partial \theta}$
- Adjoint Sensitivity method
 - Differentiating Neural ODE ends up with another ODE problem.
 - Again, we use an ODE solver to compute the gradients.
 - Only memorize the final state $z(t_1)$, then $z(t)$ is exactly reproducible.

Differentiation in discrete-time
(weight-tied ResNet)

$$\frac{\partial \mathcal{L}}{\partial z_t} = \frac{\partial \mathcal{L}}{\partial z_{t+1}} + \frac{\partial \mathcal{L}}{\partial z_{t+1}} \frac{\partial f_\theta(z_t, t)}{\partial z_t}$$


$$\frac{\partial \mathcal{L}}{\partial \theta} = \sum_t \frac{\partial \mathcal{L}}{\partial z_{t+1}} \frac{\partial f_\theta(z_t, t)}{\partial \theta}$$



Differentiation in continuous-time
(Neural ODE)

$$\frac{\partial}{\partial t} \frac{\partial \mathcal{L}}{\partial z(t)} = - \frac{\partial \mathcal{L}}{\partial z(t)} \frac{\partial f_\theta(z, t)}{\partial z}$$

$$\frac{\partial \mathcal{L}}{\partial \theta} = \int_{t_0}^{t_1} \frac{\partial \mathcal{L}}{\partial z(t)} \frac{\partial f_\theta(z, t)}{\partial \theta} dt$$



Backward pass through Neural ODE

Takeaways:

1. Only need to record the final state $z(t_1)$.
2. Therefore, backward uses $\mathcal{O}(1)$ memory.
3. Backward solves another ODE problem.

- Adjoint sensitivities

- $\frac{\partial z(t)}{\partial t} = f_\theta(z, t)$, $z(1) = z(1)$
- $\frac{\partial a(t)}{\partial t} = -a(t) \cdot \frac{\partial f_\theta(z, t)}{\partial z}$, $a(1) = \frac{\partial \mathcal{L}}{\partial z(1)}$
- $\frac{\partial d(t)}{\partial t} = -a(t) \cdot \frac{\partial f_\theta(z, t)}{\partial \theta}$, $d(1) = 0$

$z(t)$	hidden state
$a(t) = -\frac{\partial \mathcal{L}}{\partial z(t)}$	adjoint state
$d(t) = \frac{\partial \mathcal{L}}{\partial \theta}$	parameter gradient


- Solve the augmented ODE.

- Define the augmented state $s(t) = \begin{bmatrix} z(t) \\ a(t) \\ d(t) \end{bmatrix}$

$$\frac{\partial s(t)}{\partial t} = \begin{bmatrix} \frac{\partial z(t)}{\partial t} \\ \frac{\partial a(t)}{\partial t} \\ \frac{\partial d(t)}{\partial t} \end{bmatrix} = \begin{bmatrix} f_\theta(z, t) \\ -a(t) \cdot \frac{\partial f_\theta(z, t)}{\partial z} \\ -a(t) \cdot \frac{\partial f_\theta(z, t)}{\partial \theta} \end{bmatrix}$$



$$s(1) = \begin{bmatrix} z(1) \\ a(1) \\ d(1) \end{bmatrix} = \begin{bmatrix} z_1 \\ \frac{\partial \mathcal{L}}{\partial z_1} \\ \mathbf{0} \end{bmatrix}$$

- Solve the IVP: $s(0) = \text{ODEint}\left(\frac{\partial s(t)}{\partial t}, s(1), [1, 0]\right)$




Neural ODE: Summary

- Inference and training Neural ODEs
 - Neural ODEs = depth-continuous residual networks
 - Forward pass runs an ODEsolve.
 - Backward pass is also an ODEsolve. (Adjoint sensitivity method)
- Training Neural ODEs (vs. traditional neural networks)
 - Constant memory regardless of depth of layers. (vs. $\mathcal{O}(L)$ memory)
 - No need to store the activations of each intermediate layer.
 - Gradient computation runs an ODEsolve in reversed time t .
- When to use Neural ODEs
 - Trajectory of the feature vector is important (e.g., continuous time series)
 - Using normalizing flows (easier change of variables)
- DEQ and Neural ODE
 - $\mathcal{O}(1)$ memory training.
 - Provides a mechanism to balance between numerical precision vs. latency at test-time.
 - Infinite / adjustable depth
 - Adaptive computation depending on the complexity of the problem.



Background: Normalizing Flows



Brubaker et al., Introduction to Normalizing Flows (ECCV2020 Tutorial)

• Normalizing Flow

- learns mapping from a complex distribution $p_{\mathcal{X}}$ into a simple distribution $p_{\mathcal{Z}}$.

1. Uses invertible nonlinear transformations. ("flow")

- $\mathbf{z} = f_{\theta}(\mathbf{x})$
- $\mathbf{x} = f_{\theta}^{-1}(\mathbf{z})$

*Requires special invertible architectures.

2. Change of variables formula

- $p_{\mathcal{X}}(\mathbf{x}) = p_{\mathcal{Z}}(\mathbf{z}) \cdot \left| \det \left(\frac{\partial \mathbf{z}}{\partial \mathbf{x}} \right) \right|$
- Typically, $p_{\mathcal{Z}}$ is fixed to a unit gaussian distribution.
- Gives ability to compute $p(\mathbf{x})$ directly.

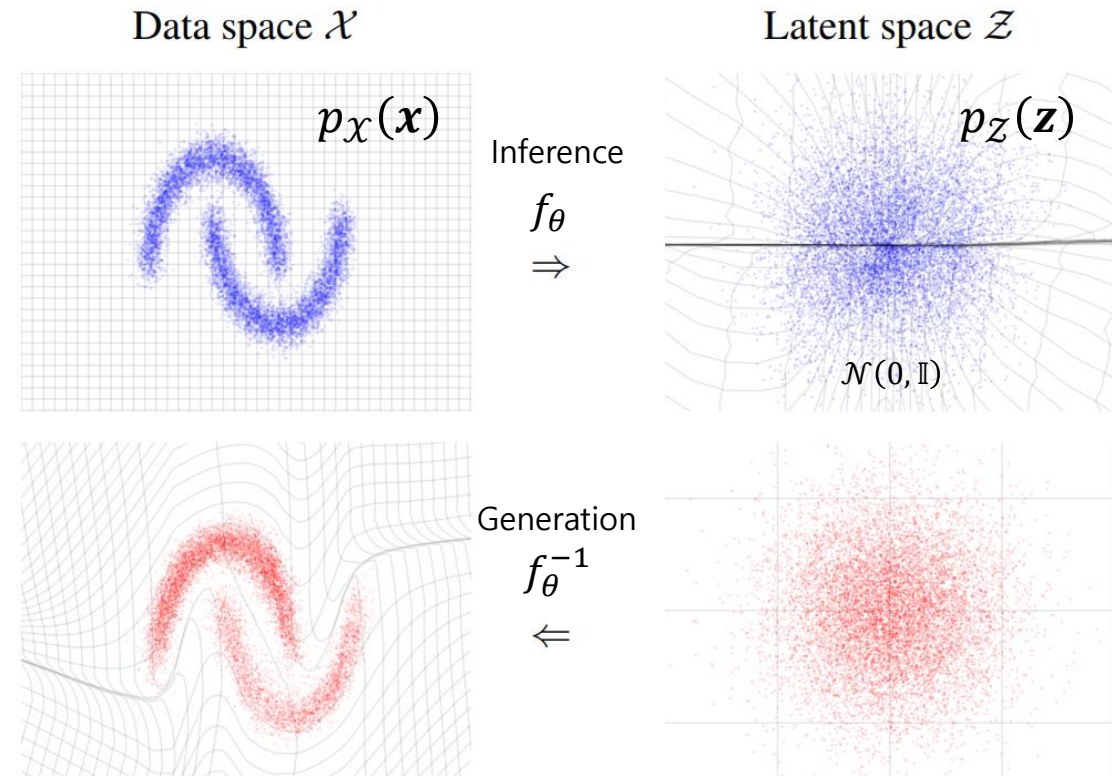
*Requires Jacobians to be structured.

3. Increases the flexibility by stacking layers of flow.

- $f_{\theta} = f_K \circ \dots \circ f_2 \circ f_1$

4. Training normalizing flows for density estimation

- $\theta^* = \underset{\theta}{\operatorname{argmax}} \log p(\mathbf{x}) = \underset{\theta}{\operatorname{argmax}} \left\{ \log p(f_{\theta}(\mathbf{x})) + \log \left| \det \left(\frac{\partial f_{\theta}(\mathbf{x})}{\partial \mathbf{x}} \right) \right| \right\}$




Density Estimation using Real NVP. Dinh, Sohl-Dickstein, Bengio (2017)

Background: Normalizing Flows


- Why should you care about NF?
 - Can learn tractable probabilistic density model of data $p(\mathbf{x})$.
 - Can fit complex posterior distributions.
 - Useful for density estimation, variational inference, generative models.

Learning flexible variational distributions



Rezende et al., Variational Inference with Normalizing Flows. ICML 2015

Generating from high-dimensional data distributions




Kingma et al., Glow: Generative Flow with Invertible 1×1 Convolutions. NIPS 2018.

Continuous Normalizing Flows

Normalizing Flow


- $z_K = f_K \circ \dots \circ f_2 \circ f_1(z_0)$
- Evaluating the density
 - Change of variables
 - $\log p(z_K) = \log p(z_0) - \sum \log \left| \det \left(\frac{\partial z_{i+1}}{\partial z_i} \right) \right|$
- Challenges with NF
 - Inverting transformation is expensive – $\mathcal{O}(D^3)$
 - Jacobian determinant is expensive – $\mathcal{O}(D^3)$
 - Consequently, flow layers must be limited to have structured Jacobians.



Continuous Normalizing Flow

- $\frac{\partial z(t)}{\partial t} = f_\theta(z, t)$
- Evaluating the density
 - Instantaneous change of variables
 - $\frac{\partial \log p(z)}{\partial t} = -\text{tr} \left(\frac{\partial f(z, t)}{\partial z} \right)$
 - $\log p(z) = \log p(z_0) - \int_0^T \text{tr} \left(\frac{\partial f(z, t)}{\partial z} \right) dt$
- Advantages of CNF
 - Inverting is cheap. (solving reversed dynamics)
 - Trace is cheap – $\mathcal{O}(D)$ cost
 - Free-form Jacobian.
- Hutchinson trace estimator
 - $\text{tr}(A) = \mathbb{E}_{v \sim \mathcal{N}(0,1)} [v^\top A v]$
 - $\int_0^T \text{tr} \left(\frac{\partial f_\theta(z, t)}{\partial z} \right) dt = \mathbb{E}_{v \sim \mathcal{N}(0,1)} \left[\int_0^T v^\top \frac{\partial f(z, t)}{\partial z} v dt \right]$

VJP is cheaper than
Jacobian products.



Continuous Normalizing Flows

Takeaways:


1. Neural ODE greatly reduces computational cost of NF.
2. More flexible flows by removing architectural restriction.

- Continuous Normalizing Flows
 - Tractable probability with $\mathcal{O}(D)$ change of variables
 - Able to scale normalizing flow models
 - Recent score-based training algorithms scale to 1024x1024



Conditional inpainting and colorization without retraining.



Requires iterative sampling procedure.



[Song, Sohl-Dickstein, Kingma, Abhishek, Ermon, Poole. Score-Based Generative Modeling through Stochastic Differential Equations, 2020]

Neural ODEs: Applications

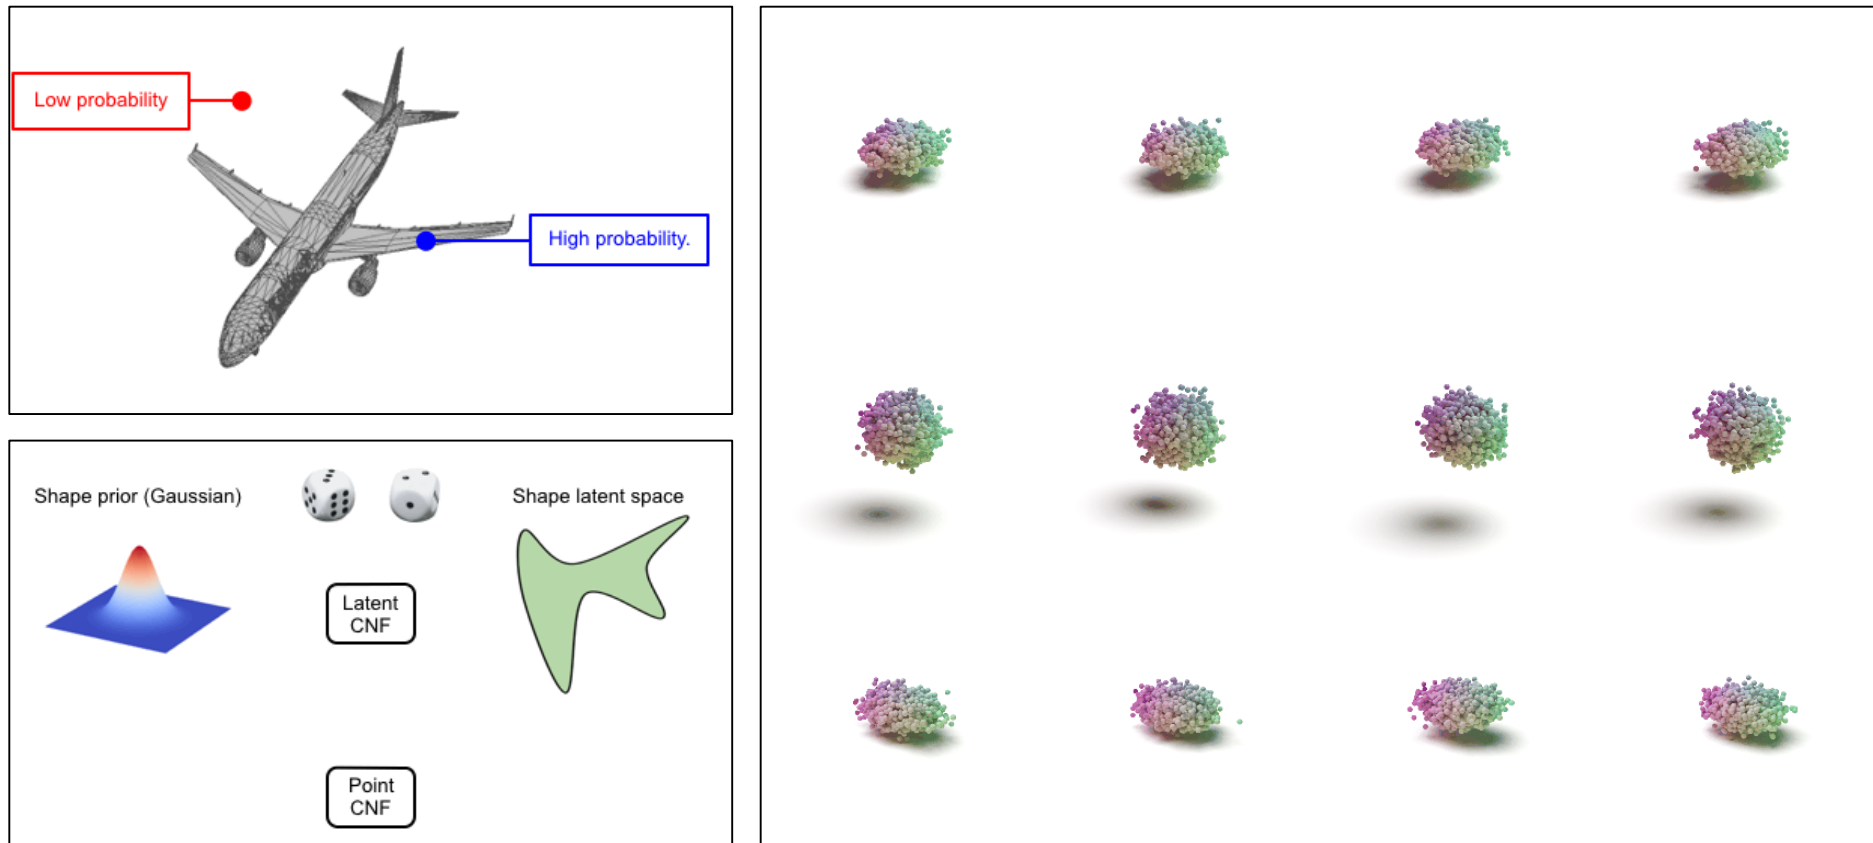
- Continuous time-series modeling
 - Can learn from datapoints sampled with irregular intervals.
- Meta-Learning
 - Rajeswaran et al., Meta-Learning with Implicit Gradients. NIPS 2019



Latent ODEs for Irregularly-Sampled Time Series. Rubanova, Chen, Duvenaud (2020)
Neural Controlled Differential Equations for Irregular Time Series.
Kidger, Morrill, Foster, Lyons (2020)
GRU-ODE-Bayes: Continuous modeling of sporadically-observed time series. de
Brouwer, Simm, Arany, Moreau. (2020)


Neural ODEs: Applications

- Learning non-intersecting transformations (homeomorphisms) using normalizing flows



Neural ODEs: Summary

- Computational advantages
 - Constant memory cost in backwards
 - Adaptive computation (trade off speed-precision flexibly)
- Modeling advantages
 - Tractable probabilistic generative models
 - Time-series models for irregularly-sampled data
 - Learning smooth homeomorphisms
- Computational disadvantages
 - Speed (future works → regularizing ODEs to be easier to solve)



Neural ODE adapts computation time according to complexity of the dynamics.

When to use DEQs vs. Neural ODEs


- Use DEQs for:
 - Drop-in implicit replacement for deep models.
 - Supervised learning (CNNs, Transformers)
 - Unsupervised learning (language modeling)
 - Shown success in large-scale tasks.
- Use Neural ODEs for:
 - Continuous-time series modeling
 - Flexible density modeling
 - Modeling homeomorphism

Simulate infinite depth with constant memory footprint

What's beyond

Open problems and future directions



- Regularizing DEQs and Neural ODEs to be faster to solve.
- Re-architecting the models to take advantage of memory efficiency.
- Scaling and application of latent stochastic differential equations.
- Partial differential equation (PDE) solution as a layer
- Differentiable optimization problem as a layer
 - Wang and Kolter et al. SATNet: Bridging deep learning and logical reasoning using a differentiable satisfiability solver. ArXiv 1905.12149
 - Diamond and Boyd. CVXPY: A Python-embedded modeling language for convex optimization. JMLR Vol. 17 (2016)
 - Agarwal and Kolter et al. Differentiable convex optimization layers. NIPS 2019.



<https://github.com/cvxgrp/cvxpy>

Regularizing implicit models to be easy to solve

- Controlling flexibility vs. speed of implicit models
 - Trade model quality for number of function evaluations (NFEs).
 - How to regularize implicit models?
- Idea so far for ODEs: regularize the dynamics to have small magnitudes.



Learning Differential Equations that are Easy to Solve. Kelly, Bettencourt, Johnson, Duvenaud. (NeurIPS 2020)

How to Train Your Neural ODE: the World of Jacobian and Kinetic Regularization. Finlay, Jacobsen, Nurbekyan, Oberman. (ICML 2020)

References

- NeurIPS 2020 Tutorial on Deep Implicit Layers: <http://implicit-layers-tutorial.org/>
- Deep Equilibrium Models
 - Shaojie Bai, Zico Kolter, Vladlen Koltun. Deep Equilibrium Models. NeurIPS 2019.
 - Shaojie Bai, Vladlen Koltun, Zico Kolter. Multiscale Deep Equilibrium Models. NeurIPS 2020.
- Neural ODEs
 - Ricky Chen, Yulia Rubanova, Jesse Bettencourt, David Duvenaud. Neural Ordinary Differential Equations. NeurIPS 2018.