

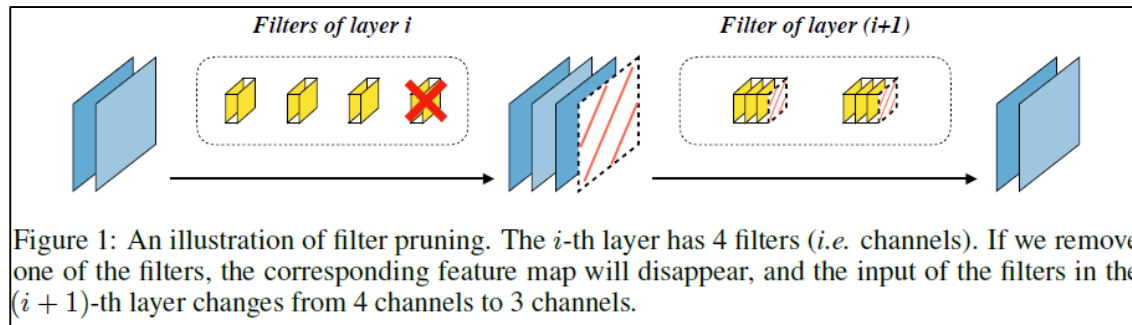
Gate Decorator:
Global Filter Pruning Method for Accelerating Deep
Convolutional Neural Networks
(NeurIPS 2019)

2020-04-13

hyounguk.shon@kaist.ac.kr

Approaches on CNN Acceleration

- Quantization
 - Reduced weight precision (Xnor-Net)
- Fast convolution
 - Factorized convolution methods (e.g., OctConv, HetConv)
- Low-rank approximation
 - Weight factorization methods (e.g., SVD)
- Filter pruning
 - Layer-by-layer pruning (e.g., ThiNet)
 - Global pruning (e.g., NISP, Gate Decorator)



Method

Problem Definition

- Global Filter Importance Ranking (GFIR)
 - The key challenge for global pruning is solving the GFIR problem.
 - Our intent is to identify the least important filter k among the model's set of filters K .
- GFIR formulated as an optimization problem:

$$k^* = \arg \min_k |\mathcal{L}(X, Y; \theta) - \mathcal{L}(X, Y; \theta_k^+)| \quad s.t. \|k\|_0 > 0$$

- i.e., search for the kernel that brings minimum change in loss when after pruned.
- Naïve search algorithm demands infeasible computational cost.

Overview of the paper

1. Gate Decorator

- Multiply each feature map by a trainable scaling factor ϕ .
- This enables an efficient importance metric for the search problem.
- Gate decorators serve for the temporary purpose of pruning.

2. Tick-Tock pruning framework

- An iterative pruning process to boost pruning accuracy.

3. Grouped pruning policy

- Resolves pruning constraints caused by shortcut connections.

- A quick look on performance:

- CIFAR-10 / ResNet-56: 70% FLOPs reduction without noticeable loss in accuracy. (SOTA)
- ImageNet / ResNet-50: 40% FLOPs reduction while *increasing* top-1 accuracy.

Gate Decorator – (I)

- Gate Decorator

1. We reparametrize a filter by introducing a gate variable ϕ .
As such, pruning a filter is equivalent to setting ϕ to zero.

$$\phi \in \mathbb{R} \text{ and use } \hat{z} = \phi z$$

$$\Delta \mathcal{L}_\Omega(\phi) = |\mathcal{L}_\Omega(\phi) - \mathcal{L}_\Omega(0)|$$

2. Assuming $|\phi| \ll 1$, we approximate $\mathcal{L}_\Omega(\phi)$ by Taylor expansion.
This approximation allows it to exploit backpropagation.

$$\begin{aligned} \mathcal{L}_\Omega(0) &= \sum_{p=0}^P \frac{\mathcal{L}_\Omega^{(p)}(\phi)}{p!} (0 - \phi)^p + R_P(\phi) \\ &= \mathcal{L}_\Omega(\phi) - \phi \nabla_\phi \mathcal{L}_\Omega + R_1(\phi) \end{aligned}$$

$$\Delta \mathcal{L}_\Omega(\phi) = |\phi \nabla_\phi \mathcal{L}_\Omega - R_1(\phi)| \approx |\phi \nabla_\phi \mathcal{L}_\Omega| = \left| \frac{\delta \mathcal{L}}{\delta \phi} \phi \right|$$

3. We evaluate important scores on the filters over the dataset D .

$$\Theta(\phi_i) = \sum_{(X,Y) \in \mathcal{D}} \left| \frac{\delta \mathcal{L}(X,Y;\theta)}{\delta \phi_i} \phi_i \right|$$

- Before pruning, we reparametrize convolution/batchnorm.
After pruning we turn them back to vanilla layers.

Gate Decorator – (II)

- Gated Convolution (GC)

1. Initialize ϕ, W .
2. Reparametrize the convolution layer.
3. Merge ϕ into weights after pruning.

$$\phi := \frac{\|W\|_F}{ck^2} \quad Y = \phi(X \otimes W) \quad W := \phi W$$

$$W := \frac{W}{\phi}$$

- Gated Batch Normalization (GBN)

- If a conv layer is followed by a BN, gate decorator is applied to BN instead.

1. Initialize $\vec{\phi}, \beta, \gamma$.
2. Reparametrize the BN layer.
3. Merge $\vec{\phi}$ into weights after pruning.

$$\vec{\phi} := \gamma$$

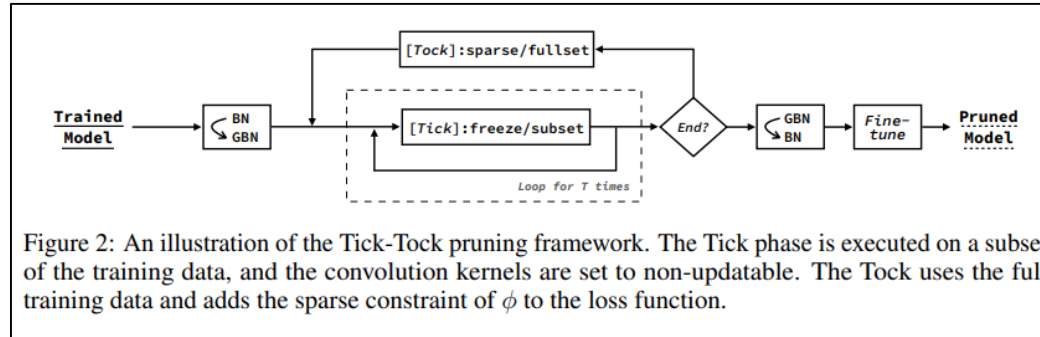
$$\beta := \frac{\beta}{\gamma} \quad \hat{z} = \frac{z_{in} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}; \quad z_{out} = \vec{\phi}(\gamma \hat{z} + \beta), \quad \vec{\phi} \in \mathbb{R}^c$$

$$\gamma := 1$$

$$\gamma := \vec{\phi}$$

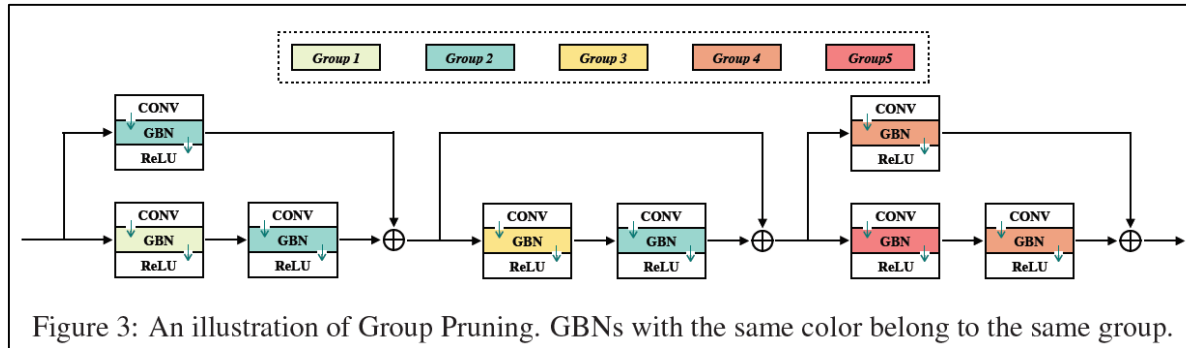
$$\beta := \beta \vec{\phi}$$

Tick-Tock pruning framework



- “Tick” phase
 - Calculates importance scores, prunes filters, fixes the internal covariate shift.
 - 1. Train the network on a small subset of the dataset for just one epoch.
We update only ϕ and the output layer to avoid overfitting.
 - 2. Compute importance score Θ , and remove a portion of the least important filters.
 - 3. Repeat the process for T times.
- “Tock” phase
 - Fine-tune the network to reduce the accumulation of errors caused by pruning.
 - Sparsity constraint on ϕ is included to training objective.
$$\mathcal{L}_{tock} = \mathcal{L} + \lambda \sum_{\phi \in \Phi} |\phi|$$
- Fine-tuning to the training set.
 - Fine-tuning trains more epochs than a Tock step.
 - Fine-tune does not include the sparsity constraint to the loss function.

Group Pruning policy



- The constrained pruning problem
 - Pruning filters in a residual block may result in misaligned feature maps.
 - Earlier approaches include:
 - bypassing such layers and only prune internal layers of residual blocks. (limits pruning ratio)
 - inserting a sampler before the first conv layer in each res block and leave the last conv layer unpruned. (add new structures to the network)
- Grouped Pruning policy
 - The idea is to prune a group of constrained filters at a time.
 - Importance score of a group is the sum of individuals.

$$\Theta(\phi_j^G) = \sum_{g \in G} \Theta(\phi_j^g)$$

Experiments

Implementation Details

- Datasets
 - Classification: CIFAR-10, CIFAR-100, CUB-200, ILSVRC-12
 - Semantic Segmentation: PASCAL VOC 2011 + SBD
- Baseline Models
 - Classification: VGGNet, ResNet
 - Semantic Segmentation: FCN
- Tick-Tock settings
 - Pruning ratio per Tick step ($T=10$)
 - ResNet: prune 0.2%
 - VGG: prune 1% per Tick step
 - Subset of training set for Tick
 - ImageNet: 100 images per class
 - CIFAR and CUB: the whole training set

Overall Comparisons – (I)

- ResNet-56 on the CIFAR-10
 - Achieved SOTA pruning ratio without noticeable loss in accuracy.

Metric	Li et al. [25]	NISP [52]	DCP-A [56]	CP [15]	AMC [14]	C-SGD [6]	GBN-40	GBN-30
FLOPs ↓%	27.6	43.6	47.1	50.0	50.0	60.8	60.1	70.3
Params ↓%	13.7	42.6	70.3	-	-	-	53.5	66.7
Accuracy ↓%	-0.02	0.03	-0.01	1.00	0.90	-0.23	-0.33	0.03

Table 1: The pruning results of ResNet-56 [\[11\]](#) on CIFAR-10 [\[20\]](#). The baseline accuracy is 93.1%.


- ResNet-50 on the ILSVRC-12

Table 2: The pruning results of ResNet-50 [\[11\]](#) on the ImageNet [\[4\]](#) dataset. "P.Top-1" and "P.Top-5" denotes the top-1 and top-5 single center crop accuracy of the pruned model on the validation set. "[Top-1] ↓" and "[Top-5] ↓" denotes the decrease in accuracy of the pruned model compared to its unpruned baseline. "Global" identifies whether the method is a global filter pruning algorithm.

Method	Global	P. Top-1	[Top-1] ↓	P. Top-5	[Top-5] ↓	FLOPs ↓%	Param ↓%
ThiNet-70 [30]	✗	72.04	0.84	90.67	0.47	36.75	33.72
SFP [12]	✗	74.61	1.54	92.06	0.81	41.80	-
GBN-60	✓	76.19	-0.31	92.83	-0.16	40.54	31.83
NISP [52]	✓	-	0.89	-	-	44.01	43.82
FPGM [13]	✗	74.83	1.32	92.32	0.55	53.50	-
ThiNet-50 [30]	✗	71.01	1.87	90.02	1.12	55.76	51.56
DCP [56]	✗	74.95	1.06	92.32	0.61	55.76	51.45
GDP [26]	✓	71.89	3.24	90.71	1.59	51.30	-
GBN-50	✓	75.18	0.67	92.41	0.26	55.06	53.40

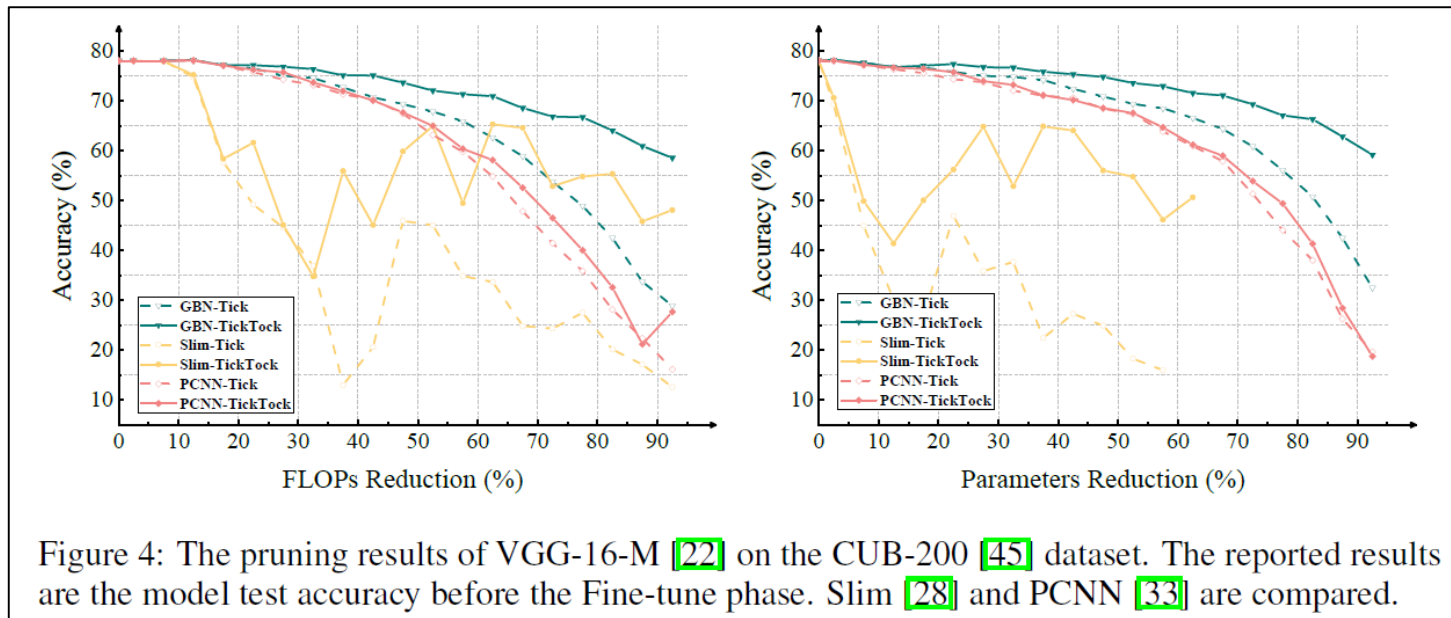
Overall Comparisons – (II)

- FCN-32s on the PASCAL VOC 2011
 - 27% FLOPs reduction, 73% parameters reduction while maintaining mIOU.




More Explorations – (I)

- Effectiveness on GFIR against other global pruning methods:
 - Slim method ranks filters by the magnitude of its factors, which is insufficient according to our analysis.
 - GBN outperforms PCNN by a large margin.



More Explorations – (II)

- Global filters pruning viewed as a NAS process:
 - Baseline model: VGG-16-M, 73.19% test acc on CIFAR-10.
 - Shrunk network: $\frac{1}{4}$ FLOPs, $\frac{1}{2}$ parameter counts, 1.98% acc drop.
 - Pruned network: $\frac{1}{4}$ FLOPs, $\sim \frac{1}{3}$ parameter counts, 1.30% acc drop.
 - Redundancy in the deep layer is unnecessary.
 - The middle layers seem to be more important.
 - This demonstrates that the pruning method can be viewed as a task-driven NAS algorithm.



More Explorations – (III)

- Effectiveness of the Tick-Tock framework:
 - Iterative pruning processes preserve better accuracy than One-shot pruning.
 - Tick-only and Tick-Tock preserves similar network structure.

FLOPs- Pruned	GBN with One-Shot			GBN with Tick-Only			GBN with Tick-Tock		
	Param	Finetune	Scratch	Param	Finetune	Scratch	Param	Finetune	Scratch
40%	79.3%	71.8	73.1	68.5%	73.0	73.7	69.0%	74.6	73.7
60%	92.0%	62.1	68.0	86.2%	71.4	72.9	85.5%	73.2	73.0
80%	97.5%	57.7	59.9	95.0%	68.4	69.6	94.7%	71.2	69.9

Table 3: The test results of VGG-16-M [22] model on the CIFAR-100 [20] dataset under different pruning schemas. The accuracy of unpruned baseline model is 73.2%. "Param" denotes the percentage of parameters that been removed. "Finetune" represents the test accuracy of the pruned model after fine-tuning. "Scratch" shows the test result of the random initialized model, which has the same architecture as the pruned one. When training the "Scratch" model, we doubled the epochs to 320.

Conclusion

- This paper proposes three components for global filter pruning:
 - The Gate Decorator algorithm to resolve the GFIR problem.
 - The Tick-Tock pruning process to boost pruning accuracy.
 - The Grouped Pruning method to resolve the constrained pruning problem.
- This paper demonstrates:
 - The global filter pruning can be viewed as a task-driven NAS algorithm.
 - Experiments show that the proposed method outperforms several SOTA pruning methods.

References

- Zhonghui You, Kun Yan, Jinmian Ye, Meng Ma, Ping Wang. Gate Decorator: Global Filter Pruning Method for Accelerating Deep Convolutional Neural Networks. *Proceedings of Advances in Neural Information Processing Systems*, 2019
- Pavlo Molchanov, Stephen Tyree, Tero Karras, Timo Aila, and Jan Kautz. Pruning convolutional neural networks for resource efficient transfer learning. ICLR, 2017.
- Zhuang Liu, Jianguo Li, Zhiqiang Shen, Gao Huang, Shoumeng Yan, and Changshui Zhang. Learning efficient convolutional networks through network slimming. ICCV, 2017.
- <https://medium.com/@nainaakash012/gate-decorator-global-filter-pruning-afc12fcc71c6>